
Disk Representations: Files, Pages, Records



Overview: Files of Pages of Records



- Overall:
 - Each table is stored in one or more OS files
 - Each file contains many pages
 - Each page contains contains many records
- Pages are the common currency understood by multiple layers:
 - Managed on disk by the disk space manager: pages read/written to physical disk/files
 - Managed in memory by the buffer manager: higher levels of DBMS only operate in memory



Files of Pages of Records



- Let's talk about a single table for now
- <u>DB FILE</u>: A collection of pages, each containing a collection of records.
- API for higher layers of the DBMS:
 - Reads:
 - Fetch a particular record by record id ...
 - Record id is a pointer encoding pair of (pageID, location on page)
 - Scan all records
 - Possibly with some conditions on the records to be retrieved
 - Updates: Insert/delete/modify record
- This abstraction could span multiple OS files and even machines

Many DB File Structures



Information is stored in files in multiple different ways

- Unordered Heap Files
 - Records placed arbitrarily across pages
- Clustered Heap Files
 - Records and pages are grouped in some meaningful way
- Sorted Files
 - Pages and records are in strict sorted order
- Index Files
 - B+ Trees, Linear Hashing, ...
 - May contain records or point to records in other files
- Focus on Unordered Heap Files for now...

Unordered Heap Files




- Collection of records in no particular order
 - Not to be confused with "heap" data-structure: efficient max/min
- As file shrinks/grows, pages (de)allocated
- To support record level operations, we must
 - Keep track of the pages in a file
 - Keep track of free space on pages
 - Keep track of the records on a page

Take 1: Heap File as List




- Heap file has one special "Header page"
 - Location of the heap file and the header page saved e.g., in catalog
- Each page contains 2 "pointers" plus free space and data
- What is wrong with this?
 - How do I find a page with enough space for a 20 byte record
 - A: Need to access many pages (w/ free space) to check



Take 2: Use a Page Directory

Berkeley CS186


- Directory, with multiple Header Pages, each encoding:
 - A pointer to page
 - #free bytes on the page
- There can be multiple such header pages
- Header pages accessed often → likely in cache
- Finding a page to fit a record required far fewer page loads than linked list. Why?
 - One header page load reveals free space of many pages
- You can imagine optimizing the page directory further
 - E.g., compressing header page, keeping header page in sorted order based on free space, etc.
 - But diminishing returns?



Summary



- Table encoded as files which are collections of pages
- Page directory provides locations of pages and free space





PAGE LAYOUT


A Note On Imagery



Data (in memory or disk) is stored in linear order



- This doesn't fit nicely on screen
 - So we will "wrap around" the linear order into a rectangle



Page Basics: The Header



- Header may contain "metadata" about the page, e.g.
 - Number of records
 - Free space
 - Maybe a next/last pointer
 - Bitmaps, Slot Table
 - (We'll talk about why all of these later)

Page Header

Things to Address



Some options:

- Record length? Fixed or Variable
- Page layout? Packed or Unpacked

Some questions:


- Find records by record id?
 - Record id = (Page, Location in Page)
- How do we add and delete records?



Fixed Length Records, Packed



- Pack records densely
- Record id = (pageId, "location in page")?
 - (pageld, record number in page)!
 - We know the offset from start of page!
 - Offset = header + (record size) x (n-1)
- Easy to add: just append
- Delete?




Fixed Length Records, Packed, Pt 2.

Record id:

(Page 2, Record 4)



- Pack records densely
- Record id = (pageId, "location in page")?
 - (pageld, record number in page)!
 - We know the offset from start of page!
- Easy to add: just append
- Delete?
 - Say we delete (Page 2, Record 3)




Fixed Length Records: Packed, Pt 3.

Record id:

(Page 2, Record 4)



- Pack records densely
- Record id = (pageId, "location in page")?
 - (pageId, record number in page)!
 - We know the offset from start of page!
- Easy to add: just append
- Delete?
 - Say we delete (Page 2, Record 3)
 - Now free space... need to reorg



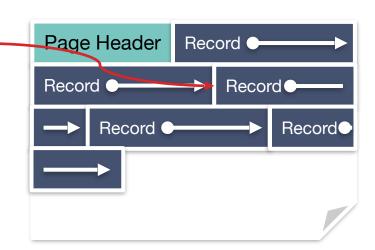
Fixed Length Records: Packed, Pt. 5

Record id:

(Page 2, Record 3)



- Pack records densely
- Record id = (pageId, "location in page")?
 - (pageld, record number in page)!
 - We know the offset from start of page!
- Easy to add: just append
- Delete?
 - Packed implies re-arrange!
 - "record id" (Page 2, Record 4) now need to be updated to (Page 2, Record 3)
 - Record Ids need to be updated!
 - Could be expensive if they're in other files.




Fixed Length Records: Unpacked



Record id: (Page 2, Record 4)

- Bitmap denotes "slots" with records
- Record id = (pageId, "location in page")?
 - (pageld, slotld)
- Insert: find first empty slot in bitmap
- Delete: ?




Fixed Length Records: Unpacked, Pt. 2 Berkeley

Record id:

(Page 2, Record 4)



- Bitmap denotes "slots" with records
- Record id = (pageId, "location in page")?
 - (pageld, slotld)
- Insert: find first empty slot in bitmap
- Delete: clear bit
 - No reorganization needed!
 - Small cost of a bitmap, which can be very compact




Variable Length Records



 We've already seen that packed isn't the best idea, so let's consider the unpacked case


- How do we know where each record begins (mapping recordid to location)?
- What happens when we add and delete records?



First: Relocate metadata to footer




We'll see why this is handy shortly...




Slotted Page

- Introduce slot directory in footer
 - Pointer to free space
 - Length + Pointer to beginning of record
 - reverse order
- Record ID = location in slot table
 - from right
- Delete?
 - e.g., 4th record on the page



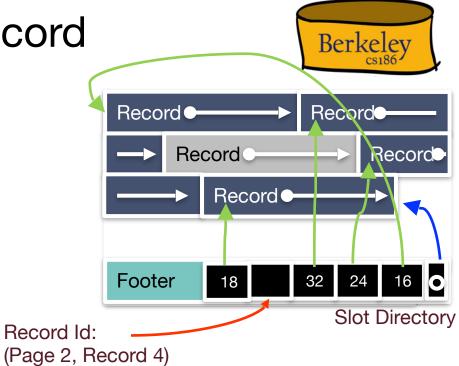
Slotted Page: Delete Record

- Delete record (Page 2, Record 4):
 - Set 4th slot directory pointer to null
 - Doesn't affect pointers to other records (no internal reorg, and no updating of external pointers)



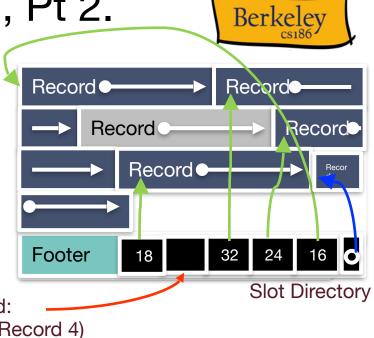
Slotted Page: Insert Record

Insert:



Slotted Page: Insert Record, Pt 2.

- Insert:
 - Place record in free space on page




Record Id:

(Page 2, Record 4)

Slotted Page: Insert Record, Pt. 3

- Insert:
 - Place record in free space on page
 - Create pointer/length pair in next open slot in slot directory




Record Id:

(Page 2, Record 4)

Slotted Page: Insert Record, Pt. 4


- Insert:
 - Place record in free space on page
 - Create pointer/length pair in next open slot in slot directory
 - Update the free space pointer
 - Fragmentation?



Record Id: (Page 2, Record 4) Slotted Page: Insert Record, Pt. 6

Berkeley CS186


- Insert:
 - Place record in free space on page
 - Create pointer/length pair in next open slot in slot directory
 - Update the free space pointer
 - Fragmentation?
 - Reorganize data on page!



(Page 2, Record 4)


Slotted Page: Leading Questions

- Reorganize data on page
 - Is this safe?
 - Yes this is safe because records ids don't change. Record ids refer to slots
- When should I reorganize?
 - We could re-organize on delete
 - Or wait until fragmentation blocks record addition and then reorganize.
 - Often pays to be a little sloppy if page never gets more records.
- What if we need more slots?
 - Let's see...



Slotted Page: Growing Slots

- Tracking number of slots in slot directory
 - Empty or full



Record Id:


(Page 2, Record 4)

Slotted Page: Growing Slots, Pt. 2

Berkeley CS186


- Tracking number of slots in slot directory
 - Empty or full
- If full slots = number of slots, then extend slot directory
- To extend slot directory
 - Slots grow from end of page inward

 - Easy!



Slotted Page: Growing Slots, Pt. 3

- Tracking number of slots in slot directory
 - Empty or full
- Extend slot directory
 - Slots grow from end of page inward
 - Records grow from beginning of page inward.
 - Easy!
- And update count



Record Id: (Page 2, Record 4)

Slotted Page: Summary

- Typically use Slotted Page
 - Good for variable and fixed length records
- Not bad for fixed length records too.
 - Why?
 - Fixed length records also have NULL fields

Record Re

Record

Record •

- NULL values can be "squashed" and indicated using a flag, avoiding full attribute length storage
- But, if we have only non-NULL fields, can be worth the optimization of fixed-length format



RECORD LAYOUT

Record Formats




- Each record in a table/relation has a fixed combo of types
- Relational databases also use same page format for data on disk or in memory
 - Save cost of conversion (known as serialization/deserialization)
- Assume System Catalog stores the Schema
 - No need to store type information with records (save space!)
 - Catalog is just another table
- Goals:
 - Fast access to fields (why?)
 - Records should be compact
- Easy Case: Fixed Length Fields
- Interesting Case: Variable Length Fields

Record Formats: Fixed Length




- Finding i'th field?
 - done via arithmetic (fast)
- Making it more compact?
 - If all fields are not-null, no good way of compacting
 - Else apply variable length techniques, next



Record Formats: Variable Length

What happens if fields are variable length?





Could store with padding? (Essentially fixed length) Wasted Space



But have to account for largest possible string (wasteful) or rearrange as soon as a larger string comes (inefficient).


Could store with delimiters (e.g., commas)?

But makes it hard to find fields and also ensure that commas are not part of the string


Record Formats: Variable Length, Pt. 7



What happens if fields are variable length?




Solution: introduce a record header



- Easy access to fields, and almost as compact as can be (modulo header)
 - Same approach can be used to squash fixed length null fields w. many nulls

Overview: Representations





Files: Summary



- DBMS "File" contains pages, and records within pages
 - Heap files: unordered records organized with directories
- Page layouts
 - Fixed-length packed and unpacked
 - Variable length records in slotted pages, with intra-page reorg
- Variable length record format
 - Direct access to i'th field and null values